

NAME: - Meetkumar Vasava

MACHINE LEARNING INTERNSHIP @ BHARAT
INTERN

PROJECT NAME - HOUSE PRICE PREDICTION

Github Link

https://github.com/MeetVasava/Bharat_intern_iris_flowe

Dataset Exploration and Preprocessing

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [3]: # Load the dataset from CSV
df = pd.read_csv('house_data.csv')
```

```
In [4]: # Exploratory Data Analysis (EDA)
# Let's take a quick look at the first few rows of the dataset
print(df.head())
```

		date	price	bedrooms	bathrooms	sqft_living	sqft_lot	\
0	2014-05-02 00:00:00	313000.0	3.0	1.50	1340	7912		
1	2014-05-02 00:00:00	2384000.0	5.0	2.50	3650	9050		
2	2014-05-02 00:00:00	342000.0	3.0	2.00	1930	11947		
3	2014-05-02 00:00:00	420000.0	3.0	2.25	2000	8030		
4	2014-05-02 00:00:00	550000.0	4.0	2.50	1940	10500		

	floors	waterfront	view	condition	sqft_above	sqft_basement	yr_built	\
0	1.5	0	0	3	1340	0	1955	
1	2.0	0	4	5	3370	280	1921	
2	1.0	0	0	4	1930	0	1966	
3	1.0	0	0	4	1000	1000	1963	
4	1.0	0	0	4	1140	800	1976	

	yr_renovated	street	city	statezip	country
0	2005	18810 Densmore Ave N	Shoreline	WA 98133	USA
1	0	709 W Blaine St	Seattle	WA 98119	USA
2	0	26206-26214 143rd Ave SE	Kent	WA 98042	USA
3	0	857 170th Pl NE	Bellevue	WA 98008	USA
4	1992	9105 170th Ave NE	Redmond	WA 98052	USA

```
In [5]: # Summary statistics of the dataset
print(df.describe())
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	\
count	4.600000e+03	4600.000000	4600.000000	4600.000000	4.600000e+03	
mean	5.519630e+05	3.400870	2.160815	2139.346957	1.485252e+04	
std	5.638347e+05	0.908848	0.783781	963.206916	3.588444e+04	
min	0.000000e+00	0.000000	0.000000	370.000000	6.380000e+02	
25%	3.228750e+05	3.000000	1.750000	1460.000000	5.000750e+03	
50%	4.609435e+05	3.000000	2.250000	1980.000000	7.683000e+03	
75%	6.549625e+05	4.000000	2.500000	2620.000000	1.100125e+04	
max	2.659000e+07	9.000000	8.000000	13540.000000	1.074218e+06	

	floors	waterfront	view	condition	sqft_above	\
count	4600.000000	4600.000000	4600.000000	4600.000000	4600.000000	
mean	1.512065	0.007174	0.240652	3.451739	1827.265435	
std	0.538288	0.084404	0.778405	0.677230	862.168977	
min	1.000000	0.000000	0.000000	1.000000	370.000000	
25%	1.000000	0.000000	0.000000	3.000000	1190.000000	
50%	1.500000	0.000000	0.000000	3.000000	1590.000000	
75%	2.000000	0.000000	0.000000	4.000000	2300.000000	
max	3.500000	1.000000	4.000000	5.000000	9410.000000	

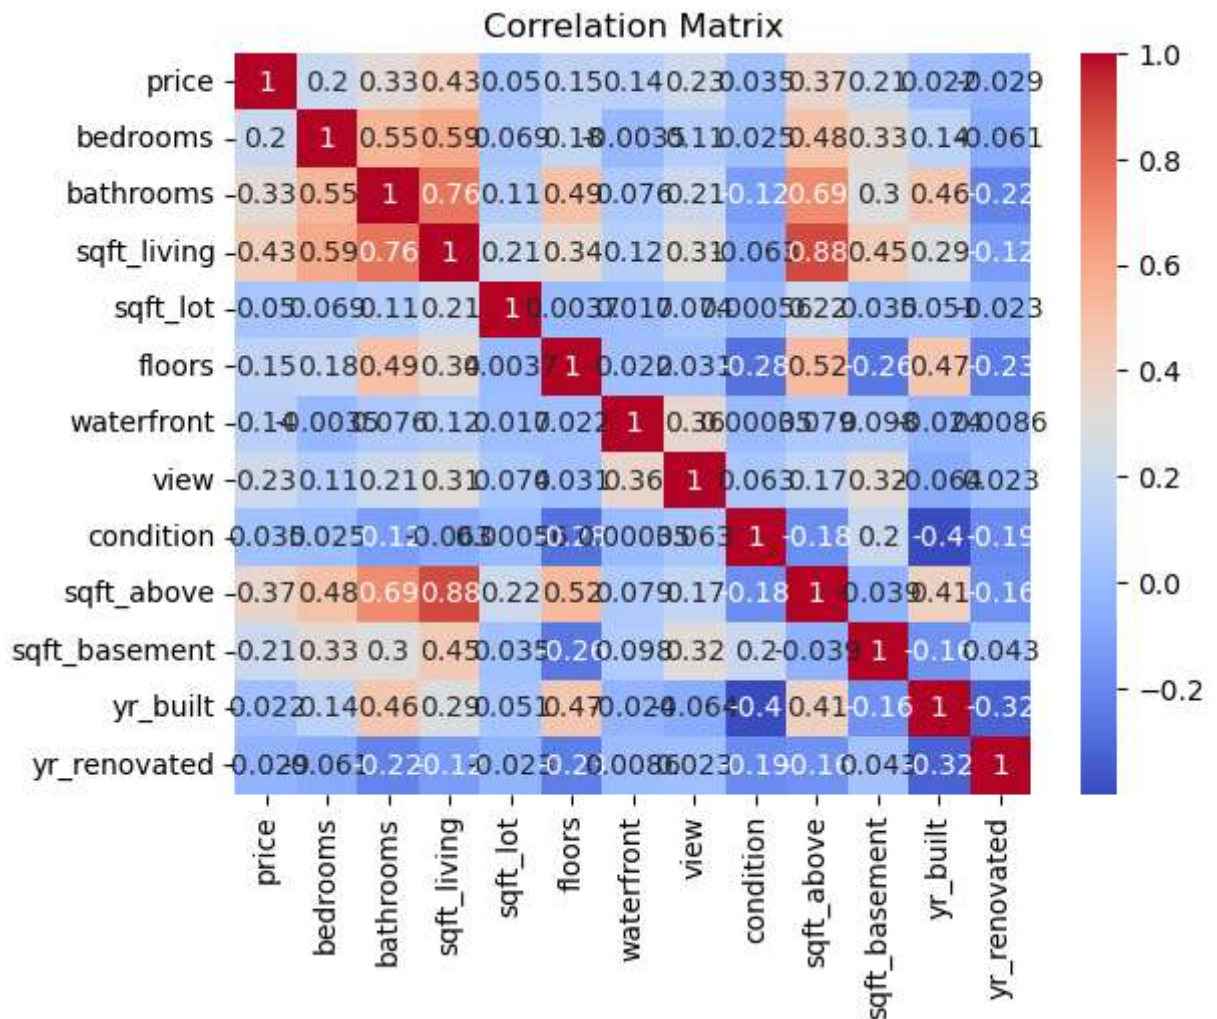
	sqft_basement	yr_built	yr_renovated
count	4600.000000	4600.000000	4600.000000
mean	312.081522	1970.786304	808.608261
std	464.137228	29.731848	979.414536
min	0.000000	1900.000000	0.000000
25%	0.000000	1951.000000	0.000000
50%	0.000000	1976.000000	0.000000
75%	610.000000	1997.000000	1999.000000
max	4820.000000	2014.000000	2014.000000

```
In [6]: # Check for missing values
print(df.isnull().sum())
```

```
date            0
price           0
bedrooms        0
bathrooms       0
sqft_living     0
sqft_lot        0
floors          0
waterfront      0
view            0
condition       0
sqft_above      0
sqft_basement   0
yr_built        0
yr_renovated    0
street          0
city            0
statezip        0
country         0
dtype: int64
```

```
In [7]: # Correlation matrix to understand feature relationships
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title("Correlation Matrix")
plt.show()
```

```
C:\Users\Meet\AppData\Local\Temp\ipykernel_3248\1819051831.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
  correlation_matrix = df.corr()
```



```
In [8]: # Preprocessing: Selecting features and target variable
X = df[['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated']]
y = df['price']

In [9]: # Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Building the Linear Regression Model

```
In [10]: # Building the Linear Regression Model
model = LinearRegression()

In [11]: # Fitting the model on the training data
model.fit(X_train, y_train)

Out[11]: ▾ LinearRegression
LinearRegression()
```

Model Evaluation

```
In [12]: # Model Evaluation
y_pred = model.predict(X_test)
```


```
In [13]: # Mean Squared Error and R-squared for model evaluation
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
In [14]: print("Mean Squared Error:", mse)
print("R-squared:", r2)
```

Mean Squared Error: 986869414953.9803
R-squared: 0.032335189956324784


Predictions and Visualization

```
In [19]: # Predictions and Visualization
# To visualize the predictions against actual prices, we'll use a scatter plot
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Prices vs. Predicted Prices")
plt.show()
```



```
In [20]: # We can also create a residual plot to check the model's performance
residuals = y_test - y_pred
plt.scatter(y_test, residuals)
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel("Actual Prices")
```

```
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.show()
```



```
In [21]: # Lastly, let's use the trained model to make predictions on new data and visualize the
new_data = [[3, 2, 1500, 4000, 1, 0, 0, 3]]
predicted_price = model.predict(new_data)

print("Predicted Price:", predicted_price[0])

Predicted Price: 331038.96876928967
C:\Users\Meet\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(
```

In []:

In []: